

DAYANANDA SAGAR COLLEGE OF ENGINEERING
(An autonomous Institution affiliated by VTU, Belagavi)
Shavige Malleshwara Hills, Kumaraswamy Layout, Bangalore-560078



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING (DATA SCIENCE)

Lab Manual
Of
Cloud Computing and
Virtualization

Semester: VI
Course Code: 22CD62
Credit: 4

Faculty In charge: Dr.Aarthi
(Associate Professor)



DAYANANDA SAGAR COLLEGE OF ENGINEERING
(An autonomous Institution affiliated by VTU, Belagavi)

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
[DATA SCIENCE]

SNO.	PROGRAM
1.	Understand the components and layers of the NIST Cloud Computing Reference Architecture.
2.	Implement and configure virtual machines using a hypervisor.
3.	Implement storage virtualization and manage virtual clusters.
4.	Introduction to Docker and its components.
5.	Deploy applications using Kubernetes for orchestration.
6.	Deploy applications on public cloud platforms (AWS, Azure).
7.	Understand IAM challenges and implementation.
8.	Identify and mitigate virtualization-specific security threats.
9.	Design a hybrid cloud architecture combining on-premises and public cloud resources.
10.	Optimize cloud infrastructure performance and monitor resource utilization.

Program 1: Understand the components and layers of the NIST Cloud Computing Reference Architecture.

Theory:

1. **NIST Cloud Model:** Defined five essential characteristics, three service models (IaaS, PaaS, SaaS), and four deployment models (Public, Private, Hybrid, Community).
2. **Key Components:** Included cloud consumer, provider, broker, auditor, and carrier, forming the cloud ecosystem.
3. **Cloud Service Interactions:** Described how different stakeholders interacted within cloud environments to ensure security, scalability, and service delivery.

Experiment:

1. **Studied** the official NIST Cloud Computing Reference Architecture document to understand its components and functions.
2. **Identified** the roles of cloud consumer, provider, broker, auditor, and carrier in a cloud environment.
3. **Explored** real-world cloud service providers (AWS, Azure, GCP) and mapped their architecture to NIST's framework.
4. **Examined** how different service models (IaaS, PaaS, SaaS) were implemented by cloud providers.
5. **Analyzed** the deployment models and observed how organizations selected public, private, hybrid, or community cloud solutions.
6. **Compared** security, governance, and compliance practices in cloud platforms with NIST's recommendations.
7. **Reviewed** case studies of companies implementing cloud solutions and aligned them with NIST's architectural principles.
8. **Documented** findings and created a summary report on how modern cloud providers align with the NIST model.

Program 2: Implement and configure virtual machines using a hypervisor.

Theory:

1. **QEMU-KVM Overview:** QEMU is a hardware emulator, and KVM (Kernel-based Virtual Machine) enables efficient virtualization on Linux.
2. **Libvirt and User Permissions:** Libvirt provides tools for managing VMs, requiring user permissions for access.
3. **Running a Virtual Machine:** After installing QEMU-KVM, an Alpine Linux ISO can be booted in a virtualized environment.

Experiment:

```
# Update the package list to ensure the latest versions are installed  
sudo apt update
```

```
# Install QEMU, KVM, and necessary virtualization tools  
sudo apt install qemu-kvm libvirt-daemon-system libvirt-clients bridge-utils virt-manager
```

```
# Add the current user to the libvirt group for VM management permissions  
sudo usermod -aG libvirt $(whoami)
```

```
# Enable and start the libvirt service to manage virtual machines  
sudo systemctl enable --now libvirtd
```

```
# Download the Alpine Linux ISO image to the Downloads directory  
wget https://dl-cdn.alpinelinux.org/alpine/v3.21/releases/x86_64/alpine-standard3.21.3-x86_64.iso -P ~/Downloads
```

```
# Start a QEMU virtual machine with the Alpine Linux ISO, booting from the CD-ROM and allocating 512MB RAM  
qemu-system-x86_64 -boot d -cdrom ~/Downloads/alpine-standard-3.21.3-x86_64.iso -m 512
```

Program 3: Implement storage virtualization and manage virtual clusters.

Theory:

1. **Storage Virtualization Overview:** QEMU-IMG is a disk image utility that allows creating, converting, and managing virtual disks for QEMU-based virtual machines.
2. **Disk Formats:** QEMU supports various disk formats, including **qcow2** (efficient, supports snapshots) and **raw** (faster but larger in size).
3. **Virtual Machine Storage Management:** A virtual disk is created using qemu-img, attached to a virtual machine, and managed using Linux utilities.

Experiment:

Storage Virtualization tool: qemu-img

Create a storage image

```
$ qemu-img create -f qcow2 virtual_disk.qcow2 40G
```

Boot the VM with the disk

```
$ qemu-system-x86_64 -boot d -cdrom ~/Downloads/alpine-standard-3.21.3x86_64.iso -m 256 -smp 2 -drive file=virtual_disk.qcow2,format=qcow2
```

Check the disc in the VM

```
# fdisk -l
```

Network Virtualization

```
# ip link add name br0 type bridge
```

```
# ip link set br0 up
```

```
# ip addr add 192.168.1.100/24 dev br0 # Adjust the IP as necessary
```

```
# mkdir /etc/qemu
```

```
# vim /etc/qemu/bridge.conf
```

-- add the contents allow
br0

- Run the VM with the network

```
# qemu-system-x86_64 -boot d -cdrom ~/Downloads/alpine-standard-  
3.21.3x86_64.iso -m 256 -smp 2 -drive file=virtual_disk.qcow2,format=qcow2 -netdev  
bridge,id=net0,br=br0 -device e1000,netdev=net0
```

Program 4: Introduction to Docker and its components.

Theory:

1. Docker is a containerization platform that allows developers to package applications with all their dependencies into a standardized unit called a container.
2. **Components of Docker:**
 - Docker Engine:** Runs and manages containers.
 - Docker Images:** Pre-configured environments for running applications.
 - Docker Containers:** Lightweight, standalone, and executable packages of software. **Docker Hub:** A repository for sharing and downloading container images.
3. Eliminates compatibility issues, enables rapid deployment and scalability.

Experiment:

Steps to Install Docker

To install Docker on Ubuntu, follow these steps:

Step 1: Update Your System

Before installing Docker, it's good practice to update the package index to make sure you have the latest versions of the repositories:

```
sudo apt update
```

Step 2: Install Required Dependencies

Next, install some dependencies that allow apt to use repositories over HTTPS:

```
sudo apt install apt-transport-https ca-certificates curl software-properties-common
```

Step 3: Add Docker's Official GPG Key

Now, you need to add Docker's official GPG key to ensure that the Docker packages are authentic:

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o  
/usr/share/keyrings/docker-archive-keyring.gpg
```

Step 4: Add Docker's Official APT Repository

Add Docker's official repository to your system's software sources:

```
echo "deb [arch=amd64 signed-by=/usr/share/keyrings/docker-archive-keyring.gpg]  
https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee
```

```
/etc/apt/sources.list.d/docker.list > /dev/null
```

Step 5: Update the Package Index Again

Now, update the package index again to include Docker's official repository:

```
sudo apt update
```

Step 6: Install Docker

Now that everything is set up, you can install Docker:

```
sudo apt install docker-ce
```

To run a simple "Hello World" Python application on Docker, you'll need to create a basic Python script and then package it into a Docker container. Here's how you can do it step by step:

Steps:

Create a directory for your Python project: Start by creating a new directory where you will store your Python script and Docker configuration files. `mkdir python-hello-world`

```
cd python-hello-world
```

Create the Python script: Inside the `python-hello-world` directory, create a Python script named `app.py` with the following content:

```
# app.py
```

```
print("Hello, World!")
```

Create a Dockerfile: In the same directory, create a file named `Dockerfile` (without an extension) to define the Docker image for your application. The Dockerfile should look like this:

```
# Use the official Python base image
```

```
FROM python:3.9-slim
```

```
# Set the working directory inside the container
```

```
WORKDIR /app
```

```
# Copy the Python script into the container COPY
```

```
app.py /app/
```



```
# Run the Python script when the container starts CMD  
["python", "app.py"]
```

Explanation of the Dockerfile:

FROM python:3.9-slim: This tells Docker to use the official Python 3.9 slim image as the base.

WORKDIR /app: Sets /app as the working directory inside the container.

COPY app.py /app/: Copies the app.py file from your local machine into the container.

CMD ["python", "app.py"]: Specifies the command to run the Python script when the container starts.

Build the Docker image: Now, you need to build the Docker image from the Dockerfile. In the terminal, run the following command from within the python-hello-world directory:

`docker build -t python-hello-world .` This will build the Docker image with the name python-hello-world.

Run the Docker container: After building the image, you can run the container with the following command:

`docker run python-hello-world` You should see the output:

Hello, World!

Stop the container: If the container doesn't stop automatically after execution (it depends on the Docker settings), you can stop it manually by running:

```
docker ps # Get the container ID  
docker stop <container_id> # Stop the container
```

Program 5: Deploy applications using Kubernetes for orchestration.

Theory:

1. Docker Desktop with Kubernetes on Windows:

Docker Desktop provides an integrated Kubernetes environment that allows developers to deploy containerized applications locally. Enabling Kubernetes via Docker Desktop offers a single-node cluster for testing and development.

2. Namespaces in Kubernetes:

Namespaces in Kubernetes allow the division of cluster resources between multiple users or projects. This helps in organizing and isolating resources logically.

3. WordPress and MySQL Deployment:

WordPress is a PHP-based CMS that typically uses MySQL as its backend database. Using Kubernetes YAML configuration files, you can deploy both services within a namespace.

Experiment:

#Download Required Files

- Download wordpress.yaml and mysql.yaml and save them in a folder (e.g., C:\Users\<YourName>\Desktop\CCV Lab\).

#Install Docker Desktop

- Download from: <https://www.docker.com/products/docker-desktop/>
- Install and restart your system if required.

#Enable Kubernetes in Docker Desktop

- Open Docker Desktop.
- Go to Settings > Kubernetes.
- Check Enable Kubernetes.
- Click Apply & Restart.
- Wait for the Kubernetes cluster to be running (check in the Docker UI).

#Open Terminal (inside Docker Desktop)

Create a Kubernetes Namespace

```
kubectl create namespace wordpress
```

```
namespace/wordpress created
```

#Apply the YAML files

```
kubecttl apply -n wordpress -f "C:\Users\<YourName>\Desktop\CCV Lab\mysql.yaml"
```

```
secret/mysql-pass created
deployment.apps/mysql created
service/mysql created
```

```
kubecttl apply -n wordpress -f "C:\Users\<YourName>\Desktop\CCV Lab\wordpress.yaml"
```

```
persistentvolumeclaim/wordpress-pvc created
deployment.apps/wordpress created
service/wordpress created
```

Check Services

```
kubecttl get svc -n wordpress
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
mysql	ClusterIP	10.103.46.218	<none>	3306/TCP	85s
wordpress	LoadBalancer	10.110.96.54	localhost	8015:30879/TCP	110s

Access WordPress

- Open your browser and go to:
<http://localhost:<NodePort>>
(Replace <NodePort> with the port shown next to WordPress service, e.g., 32750)

Program 6: Deploy applications on public cloud platforms (AWS, Azure).

Theory:

- **Virtualization:** Virtualization allows running multiple virtual systems on a single physical machine. In this program, tools like Node.js and Strapi are installed in virtual environments managed by the system or cloud provider.
- **Cloud Deployment with Strapi:** Strapi is an open-source headless CMS built on Node.js. It enables users to create and manage content APIs quickly. The application is deployed on Strapi Cloud, which provides a public hosting environment for global accessibility.
- **Networking:** Deploying to the Strapi Cloud enables external users to interact with the application over the internet using HTTP/HTTPS protocols, demonstrating basic cloud networking concepts.
- **Storage:** The application data (like content types, media, and configurations) is stored on cloud-managed persistent storage, ensuring data availability and reliability.

Experiment:

#Install Node.js, Download and install Node.js from the official link:

<https://nodejs.org/dist/v22.16.0/node-v22.16.0-x64.msi>

#Verify Node.js and NPX installation

#Open Command Prompt and run the following command to verify installation:

```
npx -v
```

#Create a new Strapi project

#Run the following command to install and create a new Strapi project:

```
npx create-strapi@latest my-strapi-project
```

Follow the command prompts and select the default options when asked.

#Create the application locally

Choose Login / Sign Up option

```
? Do you want to use the default database (sqlite) ? (Y/n) Y
? Do you want to use the default database (sqlite) ? Yes
? Start with an example structure & data? (y/N) y
? Start with an example structure & data? Yes
? Start with Typescript? (Y/n) Y
? Start with Typescript? Yes
? Install dependencies with npm? (Y/n) Y
? Install dependencies with npm? Yes
? Initialize a git repository? (Y/n) Y
? Initialize a git repository? Yes
```

#Navigate to your project folder and run the following command:

#Use the following command to start the deployment process:

```
C:\Users\Prashanth\workspace\gonagoor-ws\cclab\test-app\my-strapi-project>npm run strapi deploy
```

```
> my-strapi-project@0.1.0 strapi
> strapi deploy
```

```
? How would you like to name your project? cclab-1
? Choose your NodeJS version 20
? Choose a region for your project Asia (Southeast)
✔ Project created successfully!
🌀 Compressing project...
🌀 Project compressed successfully!
[INFO] 🚀 Uploading project...
Upload Progress | ██████████ | 100%
[SUCCESS] ✨ Upload finished!
🚀 Deploying project to production environment...
🔗 Connecting to server to get build logs
```

- Project name
- Your appropriate version of nodejs, or follow recommended option in cmd.
- Region (choose a region closest to your users)-ASIA
- Environment (e.g., production)
- Accept defaults where applicable unless otherwise instructed.

- After deployment, a public URL will be provided.
- Save this URL — it will be used for access and in the load balancer script (in the hybrid cloud experiment).

Program 7: Understand Users and IAM in a cloud setup

Theory:

1. Need for Cloud Access Management:

Organizations with multiple users accessing shared cloud infrastructure require a centralized system to manage access and prevent unauthorized usage.

2. IAM System Overview:

Identity and Access Management (IAM) in the cloud defines **resources** (e.g., VMs, GPUs, services) and governs how users interact with them.

3. Permissions and Operations:

Operations like create, edit, or delete on resources are mapped to specific **permissions** (e.g., vm.create, gpu.use).

4. Role-Based Access Control (RBAC):

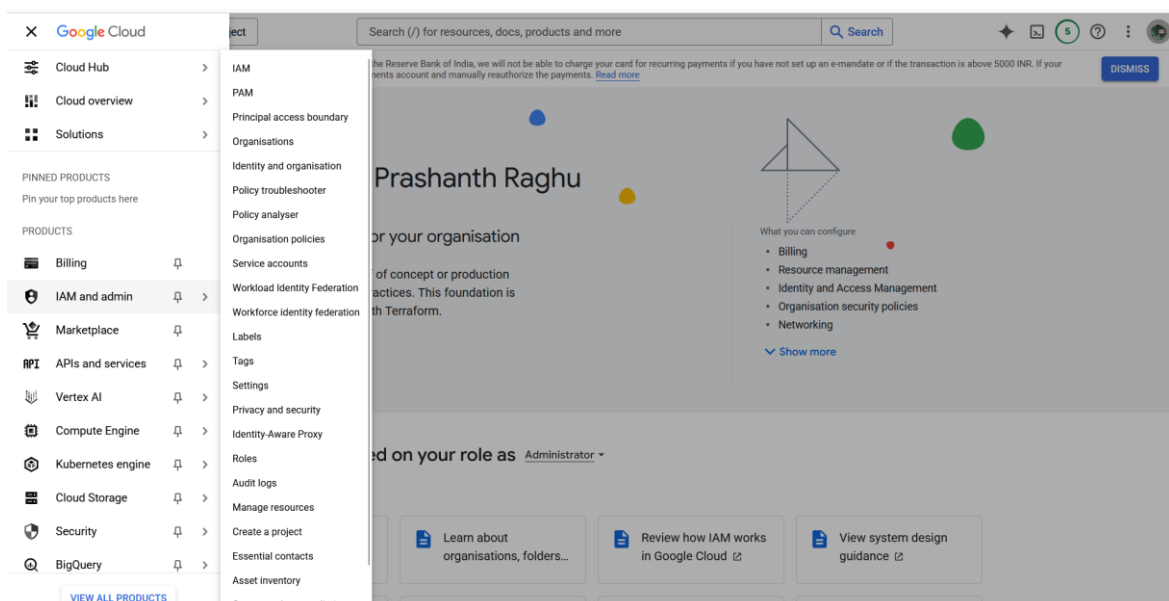
Permissions are grouped into **roles**, which act as wrappers for a set of actions. Roles simplify management by bundling access rules.

5. User Role Assignment:

Users are assigned roles based on their needs, ensuring they get the right access at the right time without overprovisioning.

Experiment:

Login to google cloud console and open the IAM Console.



Google Cloud

POC project

Search (/) for resources, docs, products and more

Q Search

5

?

:

🌐

Attention Customers from India: As of October 1st, due to regulatory changes by the Reserve Bank of India, we will not be able to charge your card for recurring payments if you have not set up an e-mandate or if the transaction is above 5000 INR. If your payments are being declined, Billing administrators will have to log into their Payments account and manually reauthorize the payments. [Read more](#)

DISMISS

IAM and admin / IAM

IAM

Learn

Allow

Deny

Recommendations history

Permissions for project POC project

These permissions affect this project and all of its resources. [Learn more](#)

☐ Include Google-provided role grants

View by principals

View by roles

Grant access

Remove access

Filter

Enter property name or value

?

⋮

Type	Principal	Name	Role	Security insights	Inher
<input type="checkbox"/>	gcp-organization-admins@gonagoortech.com		Cloud KMS Admin	Advanced security insight	g
			Organisation Administrator	Advanced security insight	g
			Pub/Sub Admin	Advanced security insight	g
			Security Centre Admin	Advanced security insight	g
<input type="checkbox"/>	prashanth@gonagoortech.com	Prashanth Raghu	Compute Admin	Advanced security insight	g
			Organisation Administrator	Advanced security insight	g
			Owner		

Recommended for you

IAM overview

Help document

Basic IAM concepts and access management in Google Cloud.

Grant an IAM role using the Google Cloud console

Help document

Use the Google Cloud console to grant IAM roles to principals at the project level

Choose predefined roles

Help document

Choose which predefined roles to grant to users.

Basic and predefined roles reference

Help document

IAM roles that you can grant to identities to access Google Cloud resources.

Manage access to projects, folders, and organizations

Help document

Grant, change, and revoke access to projects, folders, and organizations.

Troubleshoot IAM permissions

Check the available roles in the organization

Google Cloud

POC project

Search (/) for resources, docs, products and more

Q Search

5

?

:

🌐

Attention Customers from India: As of October 1st, due to regulatory changes by the Reserve Bank of India, we will not be able to charge your card for recurring payments if you have not set up an e-mandate or if the transaction is above 5000 INR. If your payments are being declined, Billing administrators will have to log into their Payments account and manually reauthorize the payments. [Read more](#)

DISMISS

IAM and admin / Roles

Roles

Create role

Create role from selection

Disable

Delete

Show info panel

Learn

Roles for POC project project

A role is a group of permissions that you can assign to principals. You can create a role and add permissions to it or copy an existing role and adjust its permissions. [Learn more](#)

Filter

Enter property name or value

?

⋮

Type	Title	Used in	Status
<input type="checkbox"/>	[Deprecated] Kubernetes Engine Node Service Agent	Service agents	Enabled
<input type="checkbox"/>	Access Approval Approver	Access Approval	Enabled
<input type="checkbox"/>	Access Approval Config Editor	Access Approval	Enabled
<input type="checkbox"/>	Access Approval Invalidator	Access Approval	Enabled
<input type="checkbox"/>	Access Approval Viewer	Access Approval	Enabled
<input type="checkbox"/>	Access Context Manager Admin	Access Context Manager	Enabled
<input type="checkbox"/>	Access Context Manager Editor	Access Context Manager	Enabled
<input type="checkbox"/>	Access Context Manager Reader	Access Context Manager	Enabled
<input type="checkbox"/>	Access Transparency Admin	Organisation Policy	Enabled
<input type="checkbox"/>	Actions Admin	Actions	Enabled
<input type="checkbox"/>	Actions Viewer	Actions	Enabled
<input type="checkbox"/>	Activity Analysis Viewer	Other	Enabled
<input type="checkbox"/>	Admin of Tenancy Units	Service Consumer Management	Enabled
<input type="checkbox"/>	Advisory Notifications admin	Other	Enabled

Workload Identity Federation

Workforce identity federation

Labels

Tags

Settings

Privacy and security

Identity-Aware Proxy

Roles

Audit logs

Essential contacts

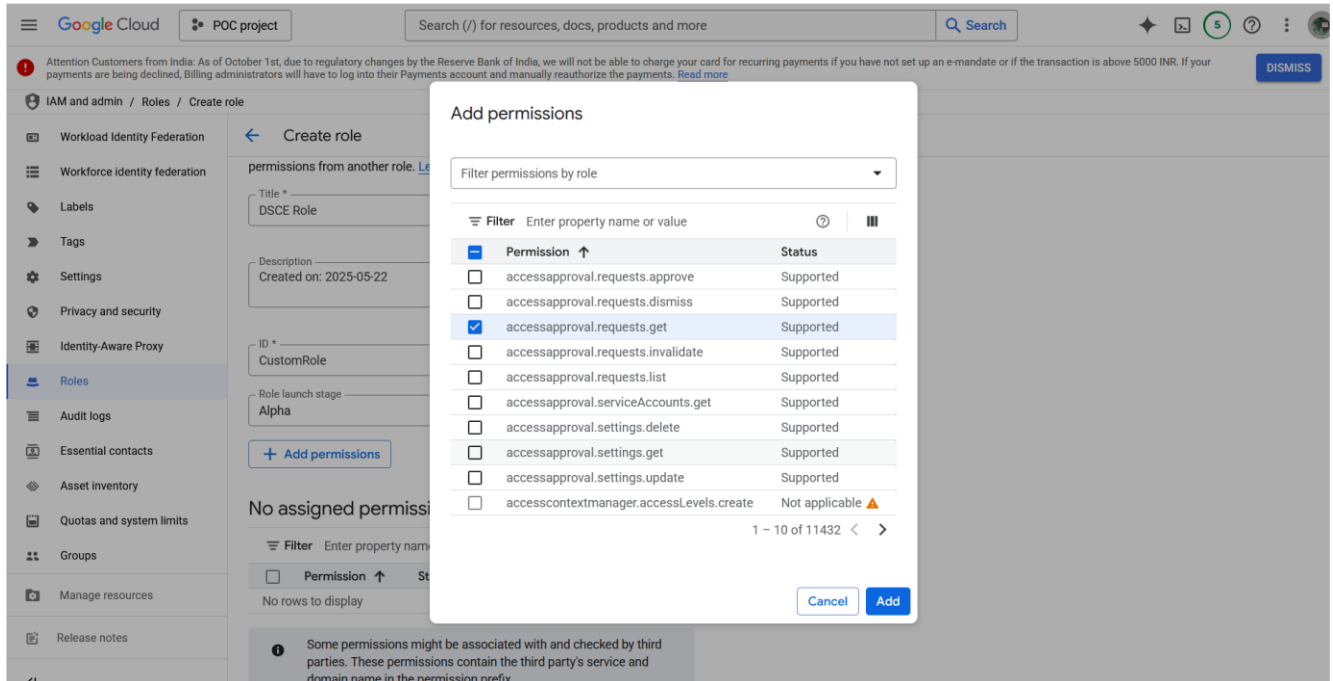
Asset inventory

Quotas and system limits

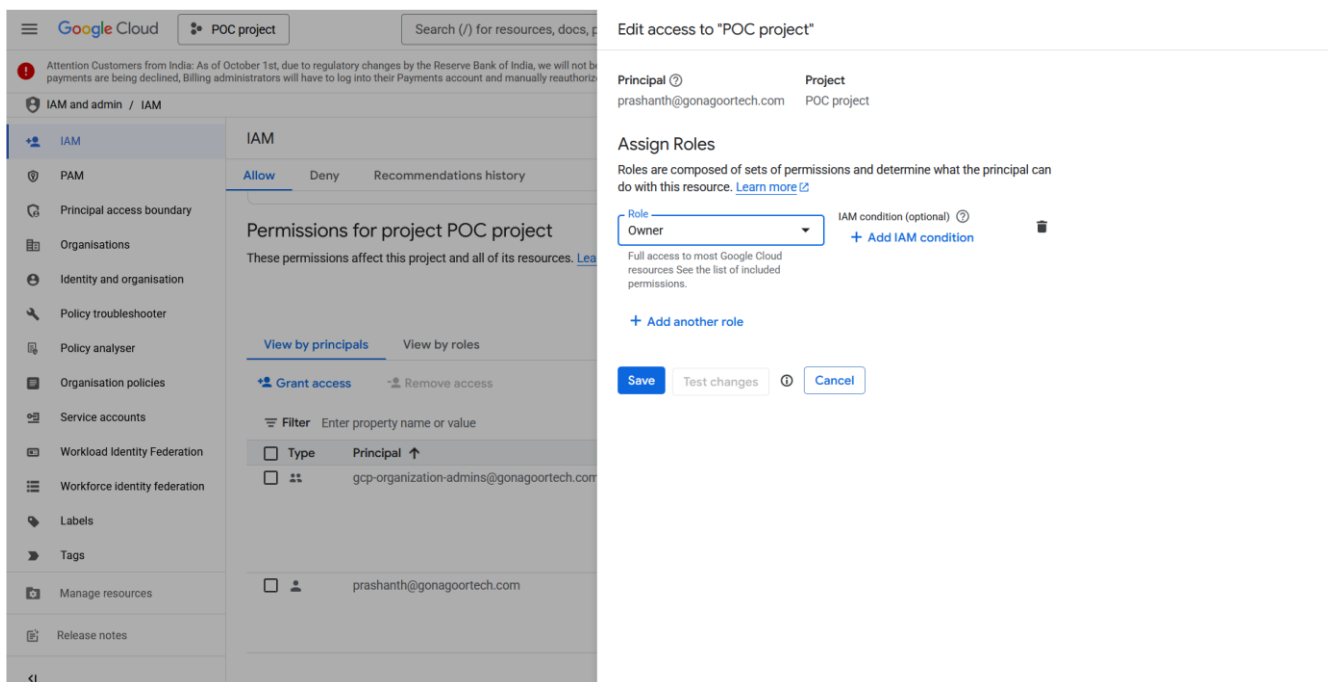
Groups

Manage resources

Release notes



Add roles to a user



Google Cloud

POC project

Search (/) for resources, docs, p

Attention Customers from India: As of October 1st, due to regulatory changes by the Reserve Bank of India, we will not b payments are being declined. Billing administrators will have to log into their Payments account and manually reauthoriz

IAM and admin / IAM

IAM

PAM

Principal access boundary

Organisations

Identity and organisation

Policy troubleshooter

Policy analyser

Organisation policies

Service accounts

Workload Identity Federation

Workforce identity federation

Labels

Tags

Manage resources

Release notes

<1

IAM

Allow

Deny

Recommendations history

Permissions for project POC project

These permissions affect this project and all of its resources. [Lea](#)

View by principals

View by roles

Grant access

Remove access

Filter

Enter property name or value

Type

Principal

gcp-organization-admins@gonagoortech.com

Type

Principal

prashanth@gonagoortech.com

Edit access to "POC project"

Principal

prashanth@gonagoortech.com

Project

POC project

Assign Roles

Roles are composed of sets of permissions and determine what the principal can do with this resource. [Learn more](#)

Role

Owner

Full access to most Google Cloud resources See the list of included permissions.

IAM condition (optional)

+ Add IAM condition

Role

DSCE Role

Created on: 2025-05-22

IAM condition (optional)

+ Add IAM condition

+ Add another role

Save

Test changes

Cancel

Summary of changes

Roles removed

N/A

Role added

DSCE Role

Test changes

Program 8: Identify and mitigate virtualization-specific security threats.

Theory:

1. Container Vulnerabilities:

Containers can have security flaws that, if ignored, may compromise both the container and host system. Regular scanning is essential before production use.

2. External Threats:

Risks can also come from outside the container, such as vulnerable runtimes or hypervisors, and must be addressed proactively.

3. Similarities to VMs and Languages:

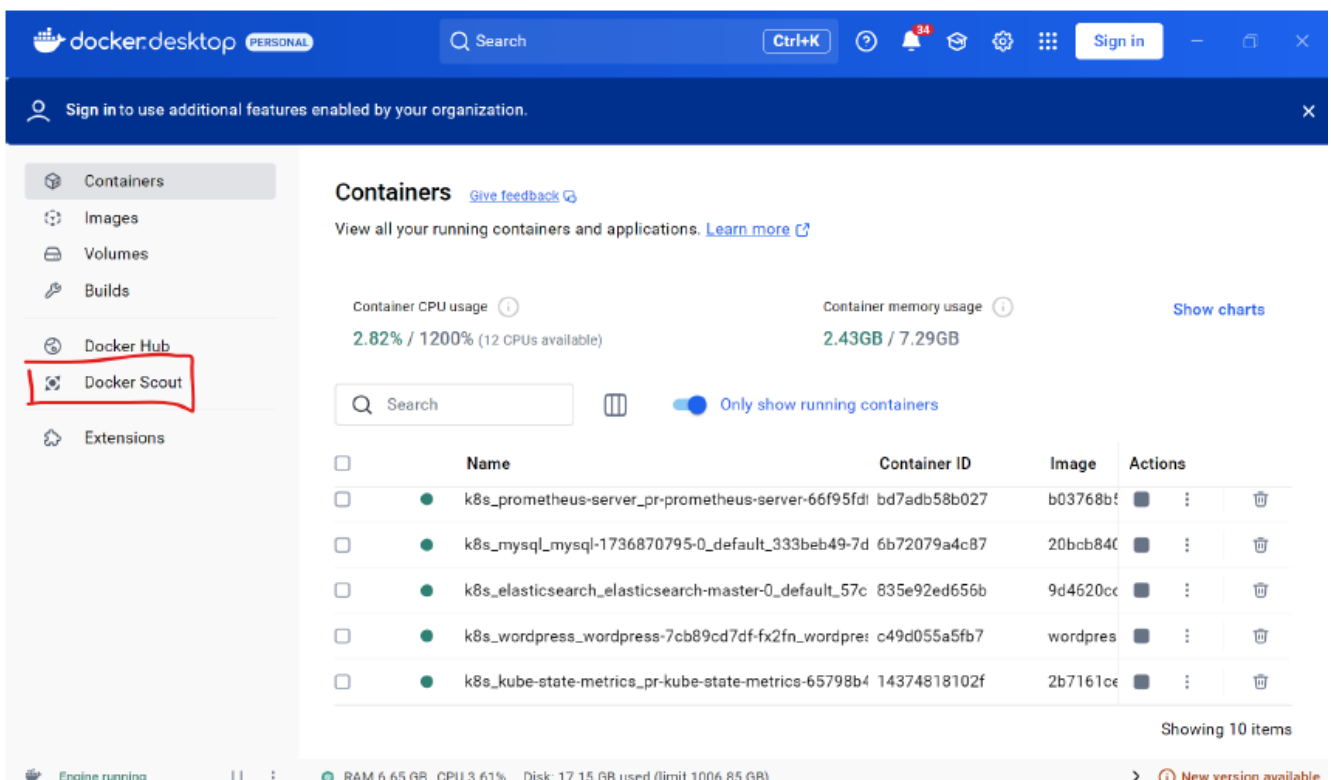
Container vulnerabilities are comparable to those in virtual machines and programming languages, requiring similar security practices.

4. Using Docker Scout:

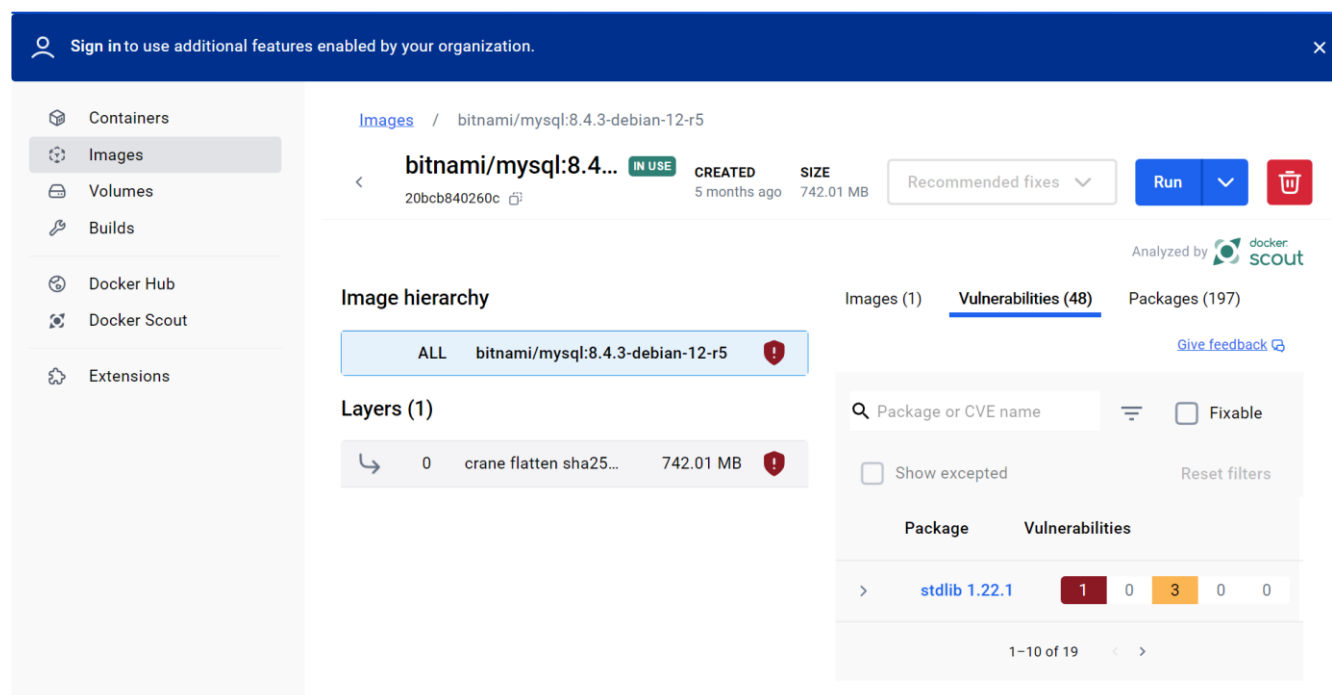
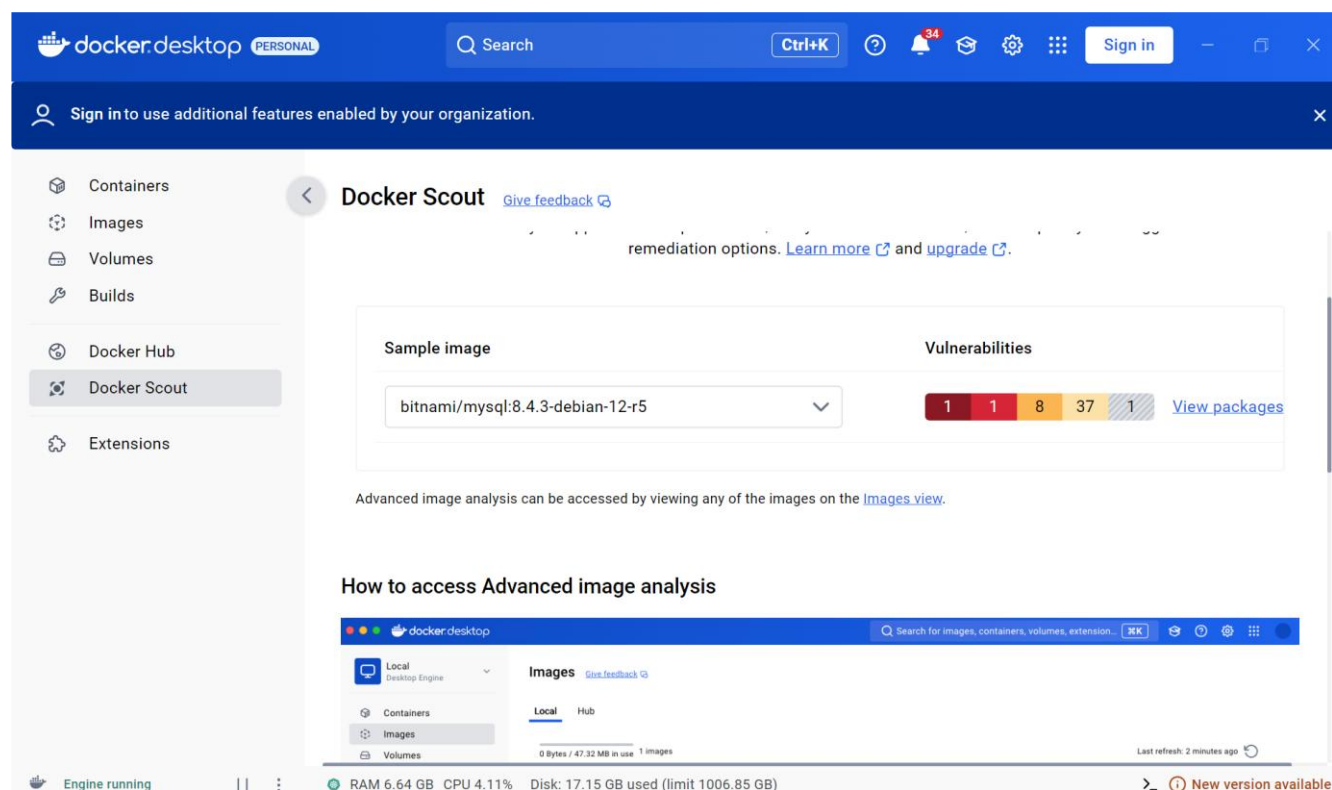
Open Docker Desktop and go to Docker Scout to scan images for known vulnerabilities before deployment.

Experiment:

Open the docker desktop console and navigate to docker scout.



Select the MySQL image and click “Check CVC” (Container Vulnerabilities Check).



Review the list of vulnerabilities.

Explore details of each vulnerability. Focus on “Critical” and “High” severity issues.

Program 9: Design a hybrid cloud architecture combining on-premises and public cloud resources.

Theory:

1. **Hybrid Cloud Architecture:** A hybrid cloud setup integrates on-premises (local) resources with public cloud services to allow flexibility, scalability, and redundancy.
2. **Use of Load Balancer:** A simple Python-based load balancer using Flask alternates traffic between a local Strapi CMS instance and a cloud-hosted Strapi application.
3. **Strapi Cloud Deployment:** The local Strapi application is built and deployed both locally and to Strapi Cloud. The cloud URL serves as the public resource, while localhost acts as the private resource.
4. **Redirection Logic:** A round-robin redirection approach is used to alternate user traffic between the local and cloud URLs using simple counter logic.

Follow the instructions as per program 6 initially and copy the application URL provided by the strapi cloud. Copy the following load balancing script as load_balancer.py

Experiment:

```
# Importing Flask web framework and redirect function
from flask import Flask, redirect

# Creating a Flask app instance
app = Flask(__name__)

# List of backend targets - local and cloud instances of Strapi CMS
TARGETS = [
    "http://localhost:1337/admin", # Localhost instance of Strapi
    "https://peaceful-confidence-a351f8d487.strapiapp.com/admin"
]

# Strapi Cloud public URL

# Counter dictionary to keep track of requests for load balancing
request_counter = {"count": 0}

# Define route for root URL
@app.route("/")
def load_balancer():
```

```

# Calculate which URL to redirect to using modulo for round-robin logic    index =
request_counter["count"] % len(TARGETS)

    target_url = TARGETS[index]

    request_counter["count"] += 1

# Redirect the user to selected backend target URL (302 by default)

    return redirect(target_url)

# Entry point of the script

if __name__ == "__main__":

    app.run(port=5000) # Run the Flask app on localhost port 5000

```

Execution Steps:

#Run your local Strapi CMS

npm run build

npm run start

#Install flask and requests libraries

pip install flask requests

#Run the Python load balancer

python load_balancer.py

#Observe the following output on the terminal

* Debug mode: off

WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.

* Running on http://127.0.0.1:5000

#Access the application in browser

Open: <http://localhost:5000/>

- First request → redirects to <http://localhost:1337/admin>
- Second request will point to public instance → redirects to <https://peaceful-confidence-a351f8d487.strapiapp.com/admin>
- Alternates with each refresh

Program 10: Optimize cloud infrastructure performance and monitor resource utilization.

Theory:

In cloud computing, optimizing infrastructure performance and monitoring resource utilization are essential to ensure efficient service delivery, cost-effectiveness, and scalability. Cloud infrastructure includes compute, storage, and network resources that support applications and services. As demand grows, proper resource allocation and monitoring become crucial to maintain system performance and availability.

1. Performance Optimization:

Cloud performance optimization involves tuning the virtualized environment to deliver high throughput, low latency, and balanced workloads. Techniques include:

- **Auto-scaling:** Automatically adding or removing resources based on demand.
- **Load Balancing:** Distributing traffic across multiple servers to prevent overload.
- **Caching:** Storing frequently accessed data in memory to reduce access time.
- **Resource Allocation:** Assigning appropriate virtual CPU, memory, and disk I/O based on workload requirements.

2. Resource Utilization Monitoring:

Monitoring involves tracking metrics such as CPU usage, memory consumption, disk I/O, and network traffic to ensure optimal operation and to detect anomalies or bottlenecks. Tools like **Prometheus**, **Grafana**, **Amazon CloudWatch**, or **Azure Monitor** provide real-time dashboards and alerts.

3. Benefits of Optimization and Monitoring:

- Improved application performance and user experience.
- Reduced operational costs through efficient resource use.
- Early detection of faults and performance degradation.
- Better planning for capacity and future scaling.

4. Key Metrics to Monitor:

- **CPU and Memory Usage:** Indicates system load and resource sufficiency.
- **Network Bandwidth:** Helps analyze data transfer efficiency.
- **Disk I/O:** Essential for applications with heavy read/write operations.
- **Latency and Response Time:** Measures service responsiveness.

By combining optimization strategies with continuous monitoring, organizations can ensure that their cloud-based systems remain resilient, responsive, and cost-effective.

