

MLOps Lab Exercises

Q1: Git Basics for ML Projects

Question: Initialize a Git repository for a new ML project. Add a Python script, configure `.gitignore` to exclude virtual environments and large files, and push to a remote repository.

Solution Guide:

Create a `Hello_world` Python file.

```
'''
```

```
git init
```

```
git add hello_world.py
```

```
git commit -m "Initial commit with hello world python script"
```

```
echo "venv/" >> .gitignore
```

```
echo "*.pyc" >> .gitignore
```

```
echo "data/*.csv" >> .gitignore
```

```
git add .gitignore
```

```
git commit -m "Add .gitignore"
```

```
git push [remote repo URL]
```

```
'''
```

Q2: Track Dataset Using DVC

Question: Initialize DVC, track a dataset file, commit changes, and push to GitHub.

Prerequisites:

- DVC installed
- A dataset file (e.g., `iris_data.csv`)

Solution Guide:

```
'''
```

```
git init
```

```
dvc init
```

```
dvc add iris_data.csv

git add iris_data.csv.dvc .gitignore .dvc/

git commit -m "Track dataset with DVC"

git push origin main

'''
```

Q3: Train a Model and Commit

****Question:**** Train a simple ML model (e.g., logistic regression) and commit the training script to Git.

- ****Solution Guide:****
 - # Save the script as train_model.py
 - from sklearn.datasets import load_iris
 - from sklearn.linear_model import LogisticRegression
 - X, y = load_iris(return_X_y=True)
 - clf = LogisticRegression(max_iter=200)
 - clf.fit(X, y)
 - print("Model trained")
 - git add train_model.py
 - git commit -m "Add model training script"

Q4: Create an Inference API (Without Docker)

****Question:**** Write a Python Flask API that serves model predictions.

- ****Solution Guide:****
 - # Save as predict_api.py
 - from flask import Flask, request, jsonify
 - import numpy as np
 - from sklearn.linear_model import LogisticRegression
 - from sklearn.datasets import load_iris
 - app = Flask(__name__)
 - X, y = load_iris(return_X_y=True)
 - model = LogisticRegression(max_iter=200).fit(X, y)
 - @app.route('/predict', methods=['POST'])
 - def predict():
 - data = request.json['features']
 - prediction = model.predict([data])
 - return jsonify({'prediction': int(prediction[0])})
 - if __name__ == '__main__':

- `app.run(debug=True)`

Q5: Dockerize Model Inference

Question: Create a Dockerfile to containerize the Flask inference API.

- **Prerequisites:**
 - `predict_api.py`
 - `requirements.txt` with Flask and scikit-learn
- **Solution Guide:**
 - # Dockerfile content
 - FROM python:3.9-slim
 - WORKDIR /app
 - COPY requirements.txt .
 - RUN pip install -r requirements.txt
 - COPY . .
 - CMD ["python", "predict_api.py"]
 - # Build and run
 - `docker build -t iris-api .`
 - `docker run -p 5000:5000 iris-api`

Q6: Streamlit UI for Model Inference

Question: Create a Streamlit app to input features and display model prediction.

- **Prerequisites:**
 - streamlit, numpy, scikit-learn installed
 - Trained model code
- **Solution Guide:**
 - # Save as `app.py`
 - import streamlit as st
 - from sklearn.datasets import load_iris
 - from sklearn.linear_model import LogisticRegression
 - import numpy as np
 - iris = load_iris()
 - X, y = iris.data, iris.target
 - model = LogisticRegression(max_iter=200).fit(X, y)
 - st.title("Iris Prediction")
 - inputs = [st.slider(label, min_value=val[0], max_value=val[1], value=val[2]) for label, val in zip(

- ['Sepal Length', 'Sepal Width', 'Petal Length', 'Petal Width'], [(4.0, 8.0, 5.1), (2.0, 4.5, 3.5), (1.0, 7.0, 1.4), (0.1, 2.5, 0.2)]]]
- if st.button('Predict'):
- result = model.predict([inputs])[0]
- st.success(f"Prediction: {iris.target_names[result]}")

Q7: Write NGINX Config for Load Balancing

Question: Write an NGINX config to load balance traffic to multiple instances of your ML API running on different ports.

- **Prerequisites:**
 - Multiple instances of model inference app running locally (e.g., ports 5000, 5001)
- **Solution Guide:**

```
# nginx.conf
http {
    upstream ml_backend {
        server localhost:5000;
        server localhost:5001;
    }
    server {
        listen 80;
        location / {
            proxy_pass http://ml_backend;
        }
    }
}
```

Q8: Simulate Traffic Using Locust

Question: Use Locust to simulate multiple users calling the model inference endpoint.

- **Prerequisites:**
 - locust installed
 - Running inference API
- **Solution Guide:**

```
# Save as locustfile.py
from locust import HttpUser, task
class MLTest(HttpUser):
    @task
    def predict(self):
```

- `self.client.post("/predict", json={"features": [5.1, 3.5, 1.4, 0.2]})`
- `# Run test`
- `locust -f locustfile.py --host=http://localhost:5000`

Q9: Bias Check & Mitigation in UCI Adult Dataset

Question: Check for bias in predictions from a model trained on UCI Adult dataset. Mitigate it and retrain.

- **Prerequisites:**
 - UCI Adult dataset
 - scikit-learn, pandas installed
- **Solution Guide:**
 - `# Load dataset and check sex-based accuracy`
 - `# Apply reweighting or oversampling to balance classes`
 - `# Retrain and compare fairness metrics`

Q10: Create a Simple GitHub Action for ML Pipeline

Question: Configure a GitHub Actions workflow to install dependencies, run `train_model.py`, and print success message.

- **Prerequisites:**
 - `train_model.py`
 - `requirements.txt`
- **Solution Guide:**
 - `# .github/workflows/train.yml`
 - `name: Train ML Model`
 - `on: [push]`
 - `jobs:`
 - `build:`
 - `runs-on: ubuntu-latest`
 - `steps:`
 - `- uses: actions/checkout@v3`
 - `- name: Set up Python`
 - `uses: actions/setup-python@v4`
 - `with:`
 - `python-version: '3.9'`
 - `- name: Install dependencies`
 - `run: pip install -r requirements.txt`
 - `- name: Run training`

- run: `python train_model.py`